



**Nelder-Mead
User's Manual
– The Fminsearch Function –**

Michaël BAUDIN

Version 0.3
September 2009

Contents

1	The <i>fminsearch</i> function	3
1.1	<i>fminsearch</i> 's algorithm	3
1.1.1	The algorithm	3
1.1.2	The initial simplex	3
1.1.3	The number of iterations	3
1.1.4	The termination criteria	4
1.2	Numerical experiments	4
1.2.1	Algorithm and numerical precision	5
1.2.2	Plot functions	9
1.2.3	Predefined plot functions	10
	Bibliography	13

Chapter 1

The *fminsearch* function

In this chapter, we analyze the implementation of the *fminsearch* which is provided in Scilab. In the first part, we describe the specific choices of this implementation with respect to the Nelder-Mead algorithm. In the second part, we present some numerical experiments which allows to check that the feature is behaving as expected, by comparison to Matlab's *fminsearch*.

1.1 *fminsearch*'s algorithm

1.1.1 The algorithm

The algorithm used is the Nelder-Mead algorithm. This corresponds to the "variable" value of the "-method" option of the *neldermead*. The "non greedy" version is used, that is, the expansion point is accepted only if it improves over the reflection point.

1.1.2 The initial simplex

The *fminsearch* algorithm uses a special initial simplex, which is an heuristic depending on the initial guess. The strategy chosen by *fminsearch* corresponds to the `-simplex0method` flag of the *neldermead* component, with the "pfeffer" method. It is associated with the `-simplex0deltausual = 0.05` and `-simplex0deltazero = 0.0075` parameters. Pfeffer's method is an heuristic which is presented in "Global Optimization Of Lennard-Jones Atomic Clusters" by Ellen Fan [2]. It is due to L. Pfeffer at Stanford. See in the help of *optimsimplex* for more details.

1.1.3 The number of iterations

In this section, we present the default values for the number of iterations in *fminsearch*.

The options input argument is an optionnal data structure which can contain the options.`MaxIter` field. It stores the maximum number of iterations. The default value is `200n`, where `n` is the number of variables. The factor 200 has not been chosen by chance, but is the result of experiments performed against quadratic functions with increasing space dimension.

This result is presented in "Effect of dimensionality on the nelder-mead simplex method" by Lixing Han and Michael Neumann [5]. This paper is based on Lixing Han's PhD, "Algorithms in Unconstrained Optimization" [4]. The study is based on numerical experiment with a quadratic function where the number of terms depends on the dimension of the space (i.e. the number

of variables). Their study shows that the number of iterations required to reach the tolerance criteria is roughly $100n$. Most iterations are based on inside contractions. Since each step of the Nelder-Mead algorithm only require one or two function evaluations, the number of required function evaluations in this experiment is also roughly $100n$.

1.1.4 The termination criteria

The algorithm used by *fminsearch* uses a particular termination criteria, based both on the absolute size of the simplex and the difference of the function values in the simplex. This termination criteria corresponds to the "tolssizedeltafvmethod" termination criteria of the *neldermead* component.

The size of the simplex is computed with the $\sigma - +$ method, which corresponds to the "sigmaplus" method of the *optimsimplex* component. The tolerance associated with this criteria is given by the "TolX" parameter of the *options* data structure. Its default value is $1.e-4$.

The function value difference is the difference between the highest and the lowest function value in the simplex. The tolerance associated with this criteria is given by the "TolFun" parameter of the *options* data structure. Its default value is $1.e-4$.

1.2 Numerical experiments

In this section, we analyse the behaviour of Scilab's *fminsearch* function, by comparison of Matlab's *fminsearch*. We especially analyse the results of the optimization, so that we can check that the algorithm is indeed behaving the same way, even if the implementation is completely different.

We consider the unconstrained optimization problem [6]

$$\min f(\mathbf{x}) \quad (1.1)$$

where $\mathbf{x} \in \mathbb{R}^2$ and the objective function f is defined by

$$f(\mathbf{x}) = 100 * (x_2 - x_1^2)^2 + (1 - x_1)^2. \quad (1.2)$$

The initial guess is

$$\mathbf{x}^0 = (-1.2, 1.)^T, \quad (1.3)$$

where the function value is

$$f(\mathbf{x}^0) = 24.2. \quad (1.4)$$

The global solution of this problem is

$$\mathbf{x}^* = (1, 1.)^T \quad (1.5)$$

where the function value is

$$f(\mathbf{x}^*) = 0. \quad (1.6)$$

1.2.1 Algorithm and numerical precision

In this section, we are concerned by the comparison of the behavior of the two algorithms. We are going to check that the algorithms produces the same intermediate and final results. We also analyze the numerical precision of the results, by detailing the number of significant digits.

To make a more living presentation of this topic, we will include small scripts which allow to produce the output that we are going to analyze. Because of the similarity of the languages, in order to avoid confusion, we will specify, for each script, the language we use by a small comment. Scripts and outputs written in Matlab's language will begin with

```
% Matlab
% ...
```

while script written in Scilab's language will begin with

```
// Scilab
// ...
```

The following Matlab script allows to see the behaviour of Matlab's *fminsearch* function on Rosenbrock's test case.

```
% Matlab
format long
banana = @(x)100*(x(2)-x(1)^2)^2+(1-x(1))^2;
[x,fval,exitflag,output] = fminsearch(banana,[-1.2, 1])
output.message
```

When this script is launched in Matlab, the following output is produced.

```
>> % Matlab
>> format long
>> banana = @(x)100*(x(2)-x(1)^2)^2+(1-x(1))^2;
>> [x,fval] = fminsearch(banana,[-1.2, 1])
>> [x,fval,exitflag,output] = fminsearch(banana,[-1.2, 1])
x =
    1.000022021783570    1.000042219751772
fval =
    8.177661197416674e-10
exitflag =
    1
output =
    iterations: 85
    funcCount: 159
    algorithm: 'Nelder-Mead_simplex_direct_search'
    message: [1x194 char]
>> output.message
ans =
Optimization terminated:
the current x satisfies the termination criteria using
OPTIONS.TolX of 1.000000e-04
and F(X) satisfies the convergence criteria using
OPTIONS.TolFun of 1.000000e-04
```

The following Scilab script allows to solve the problem with Scilab's *fminsearch*.

```
// Scilab
format(25)
function y = banana (x)
    y = 100*(x(2)-x(1)^2)^2 + (1-x(1))^2;
endfunction
[x , fval , exitflag , output] = fminsearch ( banana , [-1.2 1] )
output.message
```

The output associated with this Scilab script is the following.

```
--> // Scilab
--> format(25)
--> function y = banana (x)
--> y = 100*(x(2)-x(1)^2)^2 + (1-x(1))^2;
--> endfunction
--> [x , fval , exitflag , output] = fminsearch ( banana , [-1.2 1] )
output =
    algorithm: "Nelder-Mead_simplex_direct_search"
    funcCount: 159
    iterations: 85
    message: [3x1 string]
exitflag =
    1.
fval =
    0.0000000008177661099387
x =
    1.0000220217835567027009    1.0000422197517710998227
```

```

-->output.message
ans =

!Optimization terminated:
!
!the current x satisfies the termination criteria using OPTIONS.TolX of 1.000000e-004
!
!and F(X) satisfies the convergence criteria using OPTIONS.TolFun of 1.000000e-004
!
```

Because the two softwares do not use the same formatting rules to produce their outputs, we must perform additionnal checking in order to check our results.

The following Scilab script displays the results with 16 significant digits.

```

// Scilab
// Print the result with 15 significant digits
mprintf ( "%15e", fval );
mprintf ( "%15e_%15e", x(1), x(2) );
```

The previous script produces the following output.

```

-->// Scilab
-->mprintf ( "%15e", fval );
8.177661099387146e-010
-->mprintf ( "%15e_%15e", x(1), x(2) );
1.000022021783557e+000 1.000042219751771e+000
```

These results are reproduced verbatim in the table 1.1.

Matlab Iterations	85	
Scilab Iterations	85	
Matlab Function Evaluations	159	
Scilab Function Evaluations	159	
Matlab \mathbf{x}^*	1.000022021783570	1.000042219751772
Scilab \mathbf{x}^*	1.000022021783557e+000	1.000042219751771e+000
Matlab $f(\mathbf{x}^*)$	8.177661197416674e-10	
Scilab $f(\mathbf{x}^*)$	8.177661099387146e-010	

Fig. 1.1 : Numerical experiment with Rosenbrock's function – Comparison of results produced by Matlab and Scilab.

We must compute the common number of significant digits in order to check the consistency of the results. The following Scilab script computes the relative error between Scilab and Matlab results.

```

// Scilab
// Compare the result
xmb = [1.000022021783570 1.000042219751772 ];
err = norm(x - xmb) / norm(xmb);
mprintf ( "Relative_Error_on_x:_%e\n", err );
fmb = 8.177661197416674e-10;
err = abs(fval - fmb) / abs(fmb);
mprintf ( "Relative_Error_on_f:_%e\n", err );
```

The previous script produces the following output.

```

// Scilab
Relative Error on x : 9.441163e-015
Relative Error on f : 1.198748e-008
```

We must take into account for the floating point implementations of both Matlab and Scilab. In both these numerical softwares, double precision floating point numbers are used, i.e. the relative precision is both these softwares is $\epsilon \approx 10^{-16}$. That implies that there are approximately 16 significant digits. Therefore, the relative error on x , which is equivalent to 15 significant digits, is acceptable. If we now consider the relative error on f , which is equivalent to only 8 significant digits, that may sound as a problem. This corresponds to the square root of the relative precision,

because $\sqrt{\epsilon} \approx 10^{-8}$. In fact, this is the best that we can expect from an optimization algorithm ([1, 3]).

Therefore, the result is as close as possible to the result produced by Matlab. More specifically :

- the optimum x is the same up to 15 significant digits,
- the function value at optimum is the same up to 8 significant digits,
- the number of iterations is the same,
- the number of function evaluations is the same,
- the exit flag is the same,
- the content of the output is the same (but the string is not display the same way).

The output of the two functions is the same. We must now check that the algorithms performs the same way, that is, produces the same intermediate steps.

The following Matlab script allows to get deeper information by printing a message at each iteration with the "Display" option.

```
% Matlab
opt = optimset('Display','iter');
[x,fval,exitflag,output] = fminsearch(banana,[-1.2, 1] , opt );
```

The previous script produces the following output.

```
% Matlab
Iteration      Func-count      min f(x)      Procedure
0              1          24.2
1              3          20.05      initial simplex
2              5          5.1618      expand
3              7          4.4978      reflect
4              9          4.4978      contract outside
5             11          4.38136      contract inside
6             13          4.24527      contract inside
7             15          4.21762      reflect
8             17          4.21129      contract inside
9             19          4.13556      expand
10            21          4.13556      contract inside
11            23          4.01273      expand
12            25          3.93738      expand
13            27          3.60261      expand
14            28          3.60261      reflect
15            30          3.46622      reflect
16            32          3.21605      expand
17            34          3.16491      reflect
18            36          2.70687      expand
19            37          2.70687      reflect
20            39          2.00218      expand
21            41          2.00218      contract inside
22            43          2.00218      contract inside
23            45          1.81543      expand
24            47          1.73481      contract outside
25            49          1.31697      expand
26            50          1.31697      reflect
27            51          1.31697      reflect
28            53          1.1595      reflect
29            55          1.07674      contract inside
30            57          0.883492      reflect
31            59          0.883492      contract inside
32            61          0.669165      expand
33            63          0.669165      contract inside
34            64          0.669165      reflect
35            66          0.536729      reflect
36            68          0.536729      contract inside
37            70          0.423294      expand
38            72          0.423294      contract outside
39            74          0.398527      reflect
40            76          0.31447      expand
41            77          0.31447      reflect
42            79          0.190317      expand
43            81          0.190317      contract inside
44            82          0.190317      reflect
45            84          0.13696      reflect
46            86          0.13696      contract outside
47            88          0.113128      contract outside
```

48	90	0.11053	contract inside
49	92	0.10234	reflect
50	94	0.101184	contract inside
51	96	0.0794969	expand
52	97	0.0794969	reflect
53	98	0.0794969	reflect
54	100	0.0569294	expand
55	102	0.0569294	contract inside
56	104	0.0344855	expand
57	106	0.0179534	expand
58	108	0.0169469	contract outside
59	110	0.00401463	reflect
60	112	0.00401463	contract inside
61	113	0.00401463	reflect
62	115	0.000369954	reflect
63	117	0.000369954	contract inside
64	118	0.000369954	reflect
65	120	0.000369954	contract inside
66	122	5.90111e-005	contract outside
67	124	3.36682e-005	contract inside
68	126	3.36682e-005	contract outside
69	128	1.89159e-005	contract outside
70	130	8.46083e-006	contract inside
71	132	2.88255e-006	contract inside
72	133	2.88255e-006	reflect
73	135	7.48997e-007	contract inside
74	137	7.48997e-007	contract inside
75	139	6.20365e-007	contract inside
76	141	2.16919e-007	contract outside
77	143	1.00244e-007	contract inside
78	145	5.23487e-008	contract inside
79	147	5.03503e-008	contract inside
80	149	2.0043e-008	contract inside
81	151	1.12293e-009	contract inside
82	153	1.12293e-009	contract outside
83	155	1.12293e-009	contract inside
84	157	1.10755e-009	contract outside
85	159	8.17766e-010	contract inside

Optimization terminated:

the current x satisfies the termination criteria using OPTIONS.TolX of 1.000000e-004
and F(X) satisfies the convergence criteria using OPTIONS.TolFun of 1.000000e-004

The following Scilab script set the "Display" option to "iter" and run the *fminsearch* function.

```
// Scilab
opt = optimset ( "Display" , "iter" );
[x , fval , exitflag , output] = fminsearch ( banana , [-1.2 1] , opt );
```

```
// Scilab
Iteration   Func-count   min f(x)      Procedure
0           3          24.2          initial simplex
1           3          20.05
2           5          5.161796      expand
3           7          4.497796      reflect
4           9          4.497796      contract outside
5          11          4.3813601     contract inside
6          13          4.2452728     contract inside
7          15          4.2176247     reflect
8          17          4.2112906     contract inside
9          19          4.1355598     expand
10         21          4.1355598     contract inside
11         23          4.0127268     expand
12         25          3.9373812     expand
13         27          3.602606      expand
14         28          3.602606      reflect
15         30          3.4662211     reflect
16         32          3.2160547     expand
17         34          3.1649126     reflect
18         36          2.7068692     expand
19         37          2.7068692     reflect
20         39          2.0021824     expand
21         41          2.0021824     contract inside
22         43          2.0021824     contract inside
23         45          1.8154337     expand
24         47          1.7348144     contract outside
25         49          1.3169723     expand
26         50          1.3169723     reflect
27         51          1.3169723     reflect
28         53          1.1595038     reflect
29         55          1.0767387     contract inside
30         57          0.8834921     reflect
31         59          0.8834921     contract inside
32         61          0.6691654     expand
33         63          0.6691654     contract inside
34         64          0.6691654     reflect
35         66          0.5367289     reflect
36         68          0.5367289     contract inside
37         70          0.4232940     expand
38         72          0.4232940     contract outside
39         74          0.3985272     reflect
40         76          0.3144704     expand
41         77          0.3144704     reflect
42         79          0.1903167     expand
```


43	81	0.1903167	contract inside
44	82	0.1903167	reflect
45	84	0.1369602	reflect
46	86	0.1369602	contract outside
47	88	0.1131281	contract outside
48	90	0.1105304	contract inside
49	92	0.1023402	reflect
50	94	0.1011837	contract inside
51	96	0.0794969	expand
52	97	0.0794969	reflect
53	98	0.0794969	reflect
54	100	0.0569294	expand
55	102	0.0569294	contract inside
56	104	0.0344855	expand
57	106	0.0179534	expand
58	108	0.0169469	contract outside
59	110	0.0040146	reflect
60	112	0.0040146	contract inside
61	113	0.0040146	reflect
62	115	0.0003700	reflect
63	117	0.0003700	contract inside
64	118	0.0003700	reflect
65	120	0.0003700	contract inside
66	122	0.0000590	contract outside
67	124	0.0000337	contract inside
68	126	0.0000337	contract outside
69	128	0.0000189	contract outside
70	130	0.0000085	contract inside
71	132	0.0000029	contract inside
72	133	0.0000029	reflect
73	135	0.0000007	contract inside
74	137	0.0000007	contract inside
75	139	0.0000006	contract inside
76	141	0.0000002	contract outside
77	143	0.0000001	contract inside
78	145	5.235D-08	contract inside
79	147	5.035D-08	contract inside
80	149	2.004D-08	contract inside
81	151	1.123D-09	contract inside
82	153	1.123D-09	contract outside
83	155	1.123D-09	contract inside
84	157	1.108D-09	contract outside
85	159	8.178D-10	contract inside

Optimization terminated:
the current x satisfies the termination criteria using OPTIONS.TolX of 1.000000e-004
and F(X) satisfies the convergence criteria using OPTIONS.TolFun of 1.000000e-004

A close inspection at the data reveals that the two softwares produces indeed the same intermediate results.

1.2.2 Plot functions

In this section, we check that the plotting features of the *fminsearch* function are the same.

The following output function plots in the current graphic window the value of the current parameter **x**. To let Matlab load that script, save the content in a .m file, in a directory known by Matlab.

```
% Matlab
function stop = outfun(x, optimValues, state)
stop = false;
hold on;
plot(x(1),x(2),'.');
drawnow
```

The following Matlab script allows to perform the optimization so that the output function is called back at each iteration.

```
% Matlab
options = optimset('OutputFcn', @outfun);
[x fval] = fminsearch(banana, [-1.2, 1], options)
```

This produces the plot which is presented in figure 1.2.

The following Scilab script sets the "OutputFcn" option and then calls the *fminsearch* in order to perform the optimization.

```
// Scilab
function outfun ( x , optimValues , state )
    plot( x(1),x(2),'.');
endfunction
opt = optimset ( "OutputFcn" , outfun);
[x fval] = fminsearch ( banana , [-1.2 1] , opt );
```




Fig. 1.2 : Plot produced by Matlab's *fminsearch*, with customized output function.

The previous script produces the plot which is presented in figure 1.3.

Except for the size of the dots (which can be configured in both softwares), the graphics are exactly the same.

1.2.3 Predefined plot functions

Several pre-defined plot functions are provided with the *fminsearch* function. These functions are

- *optimplotfval*,
- *optimplotx*,
- *optimplotfunccount*.

In the following Matlab script, we use the *optimplotfval* pre-defined function.

```
% Matlab
options = optimset('PlotFcns', @optimplotfval);
[x fval] = fminsearch(banana, [-1.2, 1], options)
```

The previous script produces the plot which is presented in figure 1.4.

The following Scilab script uses the *optimplotfval* pre-defined function.

```
// Scilab
opt = optimset ( "OutputFcn" , optimplotfval );
[x fval] = fminsearch ( banana , [-1.2 1] , opt );
```

The previous script produces the plot which is presented in figure 1.5.




Fig. 1.3 : Plot produced by Scilab's *fminsearch*, with customized output function.




Fig. 1.4 : Plot produced by Matlab's *fminsearch*, with the *optimplotfval* function.

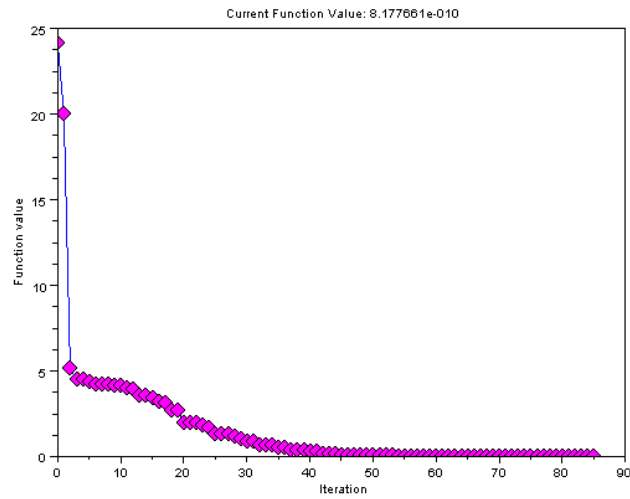


Fig. 1.5 : Plot produced by Scilab's *fminsearch*, with the *optimplotfval* function.

Bibliography

- [1] Richard P. Brent. Algorithms for minimization without derivatives, 1973.
- [2] Ellen Fan. Global optimization of lennard-jones atomic clusters. Technical report, McMaster University, February 2002.
- [3] P. E. Gill, W. Murray, and M. H. Wright. *Practical optimization*. Academic Press, London, 1981.
- [4] Lixing Han. *Algorithms in Unconstrained Optimization*. Ph.D., The University of Connecticut, 2000.
- [5] Lixing Han and Michael Neumann. Effect of dimensionality on the nelder-mead simplex method. *Optimization Methods and Software*, 21(1):1–16, 2006.
- [6] H. H. Rosenbrock. An automatic method for finding the greatest or least value of a function. *The Computer Journal*, 3(3):175–184, March 1960.