



**Nelder-Mead
User's Manual
– The Fminsearch Function –**

Michaël BAUDIN

Version 0.3
September 2009

Contents

1	The <i>fminsearch</i> function	3
1.1	<i>fminsearch</i> 's algorithm	3
1.1.1	The algorithm	3
1.1.2	The initial simplex	3
1.1.3	The number of iterations	3
1.1.4	The termination criteria	4
1.2	Numerical experiments	4
1.2.1	Algorithm and numerical precision	4
1.2.2	Plot features	9
	Bibliography	11

Chapter 1

The *fminsearch* function

In this chapter, we analyze the implementation of the *fminsearch* which is provided in Scilab. In the first part, we describe the specific choices of this implementation with respect to the Nelder-Mead algorithm. In the second part, we present some numerical experiments which allows to check that the feature is behaving as expected, by comparison to Matlab's *fminsearch*.

1.1 *fminsearch*'s algorithm

1.1.1 The algorithm

The algorithm used is the Nelder-Mead algorithm. This corresponds to the "variable" value of the "-method" option of the *neldermead*. The "non greedy" version is used, that is, the expansion point is accepted only if it improves over the reflection point.

1.1.2 The initial simplex

The *fminsearch* algorithm uses a special initial simplex, which is an heuristic depending on the initial guess. The strategy chosen by *fminsearch* corresponds to the `-simplex0method` flag of the *neldermead* component, with the "pfeffer" method. It is associated with the `-simplex0deltausual = 0.05` and `-simplex0deltazero = 0.0075` parameters. Pfeffer's method is an heuristic which is presented in "Global Optimization Of Lennard-Jones Atomic Clusters" by Ellen Fan [2]. It is due to L. Pfeffer at Stanford. See in the help of *optimsimplex* for more details.

1.1.3 The number of iterations

In this section, we present the default values for the number of iterations in *fminsearch*.

The options input argument is an optionnal data structure which can contain the options.`MaxIter` field. It stores the maximum number of iterations. The default value is `200n`, where `n` is the number of variables. The factor 200 has not been chosen by chance, but is the result of experiments performed against quadratic functions with increasing space dimension.

This result is presented in "Effect of dimensionality on the nelder-mead simplex method" by Lixing Han and Michael Neumann [5]. This paper is based on Lixing Han's PhD, "Algorithms in Unconstrained Optimization" [4]. The study is based on numerical experiment with a quadratic function where the number of terms depends on the dimension of the space (i.e. the number

of variables). Their study shows that the number of iterations required to reach the tolerance criteria is roughly $100n$. Most iterations are based on inside contractions. Since each step of the Nelder-Mead algorithm only require one or two function evaluations, the number of required function evaluations in this experiment is also roughly $100n$.

1.1.4 The termination criteria

The algorithm used by *fminsearch* uses a particular termination criteria, based both on the absolute size of the simplex and the difference of the function values in the simplex. This termination criteria corresponds to the "tolssizedeltafvmeth" termination criteria of the *neldermead* component.

The size of the simplex is computed with the $\sigma - +$ method, which corresponds to the "sigmaplus" method of the *optimsimplex* component. The tolerance associated with this criteria is given by the "TolX" parameter of the *options* data structure. Its default value is $1.e-4$.

The function value difference is the difference between the highest and the lowest function value in the simplex. The tolerance associated with this criteria is given by the "TolFun" parameter of the *options* data structure. Its default value is $1.e-4$.

1.2 Numerical experiments

In this section, we analyse the behaviour of Scilab's *fminsearch* function, by comparison of Matlab's *fminsearch*. We especially analyse the results of the optimization, so that we can check that the algorithm is indeed behaving the same way, even if the implementation is completely different.

1.2.1 Algorithm and numerical precision

We consider the unconstrained optimization problem [6]

$$\min f(\mathbf{x}) \quad (1.1)$$

where $\mathbf{x} \in \mathbb{R}^2$ and the objective function f is defined by

$$f(\mathbf{x}) = 100 * (x_2 - x_1^2)^2 + (1 - x_1)^2. \quad (1.2)$$

The initial guess is

$$\mathbf{x}^0 = (-1.2, 1.)^T, \quad (1.3)$$

where the function value is

$$f(\mathbf{x}^0) = 24.2. \quad (1.4)$$

The global solution of this problem is

$$\mathbf{x}^* = (1, 1.)^T \quad (1.5)$$

where the function value is

$$f(\mathbf{x}^*) = 0. \quad (1.6)$$

The following Matlab script allows to see the behaviour of Matlab's *fminsearch* function on Rosenbrock's test case.

```
% Matlab
format long
banana = @(x)100*(x(2)-x(1)^2)^2+(1-x(1))^2;
[x,fval,exitflag,output] = fminsearch(banana,[-1.2, 1])
output.message
```

When this script is launched in Matlab, the following output is produced.

```
>> % Matlab
>> format long
>> banana = @(x)100*(x(2)-x(1)^2)^2+(1-x(1))^2;
>> [x,fval] = fminsearch(banana,[-1.2, 1])
>> [x,fval,exitflag,output] = fminsearch(banana,[-1.2, 1])
x =
    1.000022021783570    1.000042219751772
fval =
    8.177661197416674e-10
exitflag =
    1
output =
    iterations: 85
    funcCount: 159
    algorithm: 'Nelder-Mead_simplex_direct_search'
    message: [1x194 char]
>> output.message
ans =
Optimization terminated:
the current x satisfies the termination criteria using
OPTIONS.TolX of 1.000000e-04
and F(X) satisfies the convergence criteria using
OPTIONS.TolFun of 1.000000e-04
```

The following Scilab script allows to solve the problem with Scilab's *fminsearch*.

```
// Scilab
format(25)
function y = banana (x)
    y = 100*(x(2)-x(1)^2)^2 + (1-x(1))^2;
endfunction
[x, fval, exitflag, output] = fminsearch ( banana , [-1.2 1] )
output.message
```

The output associated with this Scilab script is the following.

```
--> // Scilab
--> format(25)
--> function y = banana (x)
--> y = 100*(x(2)-x(1)^2)^2 + (1-x(1))^2;
--> endfunction
--> [x, fval, exitflag, output] = fminsearch ( banana , [-1.2 1] )
output =
    algorithm: "Nelder-Mead_simplex_direct_search"
    funcCount: 159
    iterations: 85
    message: [3x1 string]
exitflag =
    1.
fval =
    0.0000000008177661099387
x =
    1.0000220217835567027009    1.0000422197517710998227
--> output.message
ans =

!Optimization terminated:                                     !
!                                                              !
!the current x satisfies the termination criteria using OPTIONS.TolX of 1.000000e-004 !
!                                                              !
!and F(X) satisfies the convergence criteria using OPTIONS.TolFun of 1.000000e-004 !
```

Because the two softwares do not use the same formatting rules to produce their outputs, we must perform additionnal checking in order to check our results.

The following Scilab script displays the results with 16 significant digits.

```
// Scilab
// Print the result with 15 significant digits
mprintf ( "%15e", fval );
mprintf ( "%15e_%15e", x(1), x(2) );
```

The previous script produces the following output.

```
--> // Scilab
--> mprintf ( "%15e", fval );
8.177661099387146e-010
--> mprintf ( "%15e_%15e", x(1), x(2) );
1.000022021783557e+000 1.000042219751771e+000
```

Matlab Iterations	85	
Scilab Iterations	85	
Matlab Function Evaluations	159	
Scilab Function Evaluations	159	
Matlab \mathbf{x}^*	1.000022021783570	1.000042219751772
Scilab \mathbf{x}^*	1.000022021783557e+000	1.000042219751771e+000
Matlab $f(\mathbf{x}^*)$	8.177661197416674e-10	
Scilab $f(\mathbf{x}^*)$	8.177661099387146e-010	

Fig. 1.1 : Numerical experiment with Rosenbrock's function – Comparison of results produced by Matlab and Scilab.

These results are reproduced verbatim in the table 1.1.

We must compute the common number of significant digits in order to check the consistency of the results. The following Scilab script computes the relative error between Scilab and Matlab results.

```
// Scilab
// Compare the result
xmb = [1.000022021783570    1.000042219751772 ];
err = norm(x - xmb) / norm(xmb);
mprintf ( "Relative_Error_on_x:_%e\n", err );
fmb = 8.177661197416674e-10;
err = abs(fval - fmb) / abs(fmb);
mprintf ( "Relative_Error_on_f:_%e\n", err );
```

The previous script produces the following output.

```
// Scilab
Relative_Error_on_x : 9.441163e-015
Relative_Error_on_f : 1.198748e-008
```

We must take into account for the floating point implementations of both Matlab and Scilab. In both these numerical softwares, double precision floating point numbers are used, i.e. the relative precision is both these softwares is $\epsilon \approx 10^{-16}$. That implies that there are approximately 16 significant digits. Therefore, the relative error on x , which is equivalent to 15 significant digits, is acceptable. If we now consider the relative error on f , which is equivalent to only 8 significant digits, that may sound as a problem. This corresponds to the square root of the relative precision, because $\sqrt{\epsilon} \approx 10^{-8}$. In fact, this is the best that we can expect from an optimization algorithm ([1, 3]).

Therefore, the result is as close as possible to the result produced by Matlab. More specifically :

- the optimum x is the same up to 15 significant digits,
- the function value at optimum is the same up to 8 significant digits,
- the number of iterations is the same,
- the number of function evaluations is the same,
- the exit flag is the same,
- the content of the output is the same (but the string is not display the same way).

The output of the two functions is the same. We must now check that the algorithms performs the same way, that is, produces the same intermediate steps.

The following Matlab script allows to get deeper information by printing a message at each iteration with the "Display" option.

```
% Matlab
opt = optimset('Display','iter');
[x,fval,exitflag,output] = fminsearch(banana,[-1.2, 1] , opt );
```

The previous script produces the following output.

```
% Matlab
Iteration   Func-count      min f(x)      Procedure
0           1          24.2
1           3          20.05      initial simplex
2           5          5.1618      expand
3           7          4.4978      reflect
4           9          4.4978      contract outside
5          11          4.38136     contract inside
6          13          4.24527     contract inside
7          15          4.21762     reflect
8          17          4.21129     contract inside
9          19          4.13556     expand
10         21          4.13556     contract inside
11         23          4.01273     expand
12         25          3.93738     expand
13         27          3.60261     expand
14         28          3.60261     reflect
15         30          3.46622     reflect
16         32          3.21605     expand
17         34          3.16491     reflect
18         36          2.70687     expand
19         37          2.70687     reflect
20         39          2.00218     expand
21         41          2.00218     contract inside
22         43          2.00218     contract inside
23         45          1.81543     expand
24         47          1.73481     contract outside
25         49          1.31697     expand
26         50          1.31697     reflect
27         51          1.31697     reflect
28         53          1.1595      reflect
29         55          1.07674     contract inside
30         57          0.883492    reflect
31         59          0.883492    contract inside
32         61          0.669165    expand
33         63          0.669165    contract inside
34         64          0.669165    reflect
35         66          0.536729    reflect
36         68          0.536729    contract inside
37         70          0.423294    expand
38         72          0.423294    contract outside
39         74          0.398527    reflect
40         76          0.31447     expand
41         77          0.31447     reflect
42         79          0.190317    expand
43         81          0.190317    contract inside
44         82          0.190317    reflect
45         84          0.13696     reflect
46         86          0.13696     contract outside
47         88          0.113128    contract outside
48         90          0.11053     contract inside
49         92          0.10234     reflect
50         94          0.101184    contract inside
51         96          0.0794969    expand
52         97          0.0794969    reflect
53         98          0.0794969    reflect
54        100          0.0569294    expand
55        102          0.0569294    contract inside
56        104          0.0344855    expand
57        106          0.0179534    expand
58        108          0.0169469    contract outside
59        110          0.00401463   reflect
60        112          0.00401463   contract inside
61        113          0.00401463   reflect
62        115          0.000369954  reflect
63        117          0.000369954  contract inside
64        118          0.000369954  reflect
65        120          0.000369954  contract inside
66        122          5.90111e-005  contract outside
67        124          3.36682e-005  contract inside
68        126          3.36682e-005  contract outside
69        128          1.89159e-005  contract outside
70        130          8.46083e-006  contract inside
71        132          2.88255e-006  contract inside
72        133          2.88255e-006  reflect
73        135          7.48997e-007  contract inside
74        137          7.48997e-007  contract inside
75        139          6.20365e-007  contract inside
76        141          2.16919e-007  contract outside
77        143          1.00244e-007  contract inside
```

78	145	5.23487e-008	contract	inside
79	147	5.03503e-008	contract	inside
80	149	2.0043e-008	contract	inside
81	151	1.12293e-009	contract	inside
82	153	1.12293e-009	contract	outside
83	155	1.12293e-009	contract	inside
84	157	1.10755e-009	contract	outside
85	159	8.17766e-010	contract	inside

Optimization terminated:

the current x satisfies the termination criteria using OPTIONS.TolX of 1.000000e-004
and F(X) satisfies the convergence criteria using OPTIONS.TolFun of 1.000000e-004

The following Scilab script set the "Display" option to "iter" and run the *fminsearch* function.

```
// Scilab
opt = optimset ( "Display" , "iter" );
[x , fval , exitflag , output] = fminsearch ( banana , [-1.2 1] , opt );
```

```
// Scilab
Iteration      Func-count      min f(x)      Procedure
0              3              24.2          initial simplex
1              3              20.05         expand
2              5              5.161796      reflect
3              7              4.497796      contract outside
4              9              4.497796      contract inside
5             11              4.3813601     contract inside
6             13              4.2452728     reflect
7             15              4.2176247     contract inside
8             17              4.2112906     expand
9             19              4.1355598     contract inside
10            21              4.1355598     expand
11            23              4.0127268     expand
12            25              3.9373812     expand
13            27              3.602606     reflect
14            28              3.602606     reflect
15            30              3.4662211     expand
16            32              3.2160547     reflect
17            34              3.1649126     expand
18            36              2.7068692     reflect
19            37              2.7068692     expand
20            39              2.0021824     contract inside
21            41              2.0021824     contract inside
22            43              2.0021824     expand
23            45              1.8154337     contract outside
24            47              1.7348144     expand
25            49              1.3169723     reflect
26            50              1.3169723     reflect
27            51              1.3169723     reflect
28            53              1.1595038     contract inside
29            55              1.0767387     reflect
30            57              0.8834921     contract inside
31            59              0.8834921     expand
32            61              0.6691654     contract inside
33            63              0.6691654     reflect
34            64              0.6691654     reflect
35            66              0.5367289     contract inside
36            68              0.5367289     expand
37            70              0.4232940     contract outside
38            72              0.4232940     reflect
39            74              0.3985272     expand
40            76              0.3144704     reflect
41            77              0.3144704     expand
42            79              0.1903167     contract inside
43            81              0.1903167     reflect
44            82              0.1903167     reflect
45            84              0.1369602     contract outside
46            86              0.1369602     contract outside
47            88              0.1131281     contract inside
48            90              0.1105304     reflect
49            92              0.1023402     contract inside
50            94              0.1011837     expand
51            96              0.0794969     reflect
52            97              0.0794969     reflect
53            98              0.0794969     expand
54           100              0.0569294     contract inside
55           102              0.0569294     expand
56           104              0.0344855     expand
57           106              0.0179534     contract outside
58           108              0.0169469     reflect
59           110              0.0040146     contract inside
60           112              0.0040146     reflect
61           113              0.0040146     reflect
62           115              0.0003700     contract inside
63           117              0.0003700     reflect
64           118              0.0003700     contract inside
65           120              0.0003700     contract outside
66           122              0.0000590     contract inside
67           124              0.0000337     contract outside
68           126              0.0000337     contract outside
69           128              0.0000189     contract inside
70           130              0.0000085     contract inside
71           132              0.0000029     reflect
72           133              0.0000029
```





Fig. 1.2 : Plot produced by Matlab's *fminsearch*, with customized output function.

73	135	0.0000007	contract inside
74	137	0.0000007	contract inside
75	139	0.0000006	contract inside
76	141	0.0000002	contract outside
77	143	0.0000001	contract inside
78	145	5.235D-08	contract inside
79	147	5.035D-08	contract inside
80	149	2.004D-08	contract inside
81	151	1.123D-09	contract inside
82	153	1.123D-09	contract outside
83	155	1.123D-09	contract inside
84	157	1.108D-09	contract outside
85	159	8.178D-10	contract inside

Optimization terminated:
the current *x* satisfies the termination criteria using OPTIONS.TolX of 1.000000e-004
and F(X) satisfies the convergence criteria using OPTIONS.TolFun of 1.000000e-004

A close inspection at the data reveals that the two softwares produces indeed the same intermediate results.

1.2.2 Plot features

In this section, we check that the plotting features of the *fminsearch* function are the same.

The following output function plots in the current graphic window the value of the current parameter *x*. To let Matlab load that script, save the content in a .m file, in a directory known by Matlab.

```
% Matlab
function stop = outfun(x, optimValues, state)
stop = false;
hold on;
plot(x(1),x(2),'.');
drawnow
```

The following Matlab script allows to perform the optimization so that the output function is called back at each iteration.

```
% Matlab
options = optimset('OutputFcn', @outfun);
[x fval] = fminsearch(banana, [-1.2, 1], options)
```

This produces the plot which is presented in figure 1.2.




Fig. 1.3 : Plot produced by Scilab's *fminsearch*, with customized output function.

The following Scilab script sets the "OutputFcn" option and then calls the *fminsearch* in order to perform the optimization.

```
// Scilab
function outfun ( x , optimValues , state )
    plot( x(1),x(2) , ' . ' );
endfunction
opt = optimset ( "OutputFcn" , outfun );
[x fval] = fminsearch ( banana , [-1.2 1] , opt );
```

The previous script produces the plot which is presented in figure 1.3.

Except for the size of the dots (which can be configured in both softwares), the graphics are exactly the same.

Bibliography

- [1] Richard P. Brent. Algorithms for minimization without derivatives, 1973.
- [2] Ellen Fan. Global optimization of lennard-jones atomic clusters. Technical report, McMaster University, February 2002.
- [3] P. E. Gill, W. Murray, and M. H. Wright. *Practical optimization*. Academic Press, London, 1981.
- [4] Lixing Han. *Algorithms in Unconstrained Optimization*. Ph.D., The University of Connecticut, 2000.
- [5] Lixing Han and Michael Neumann. Effect of dimensionality on the nelder-mead simplex method. *Optimization Methods and Software*, 21(1):1–16, 2006.
- [6] H. H. Rosenbrock. An automatic method for finding the greatest or least value of a function. *The Computer Journal*, 3(3):175–184, March 1960.